

Quality Metrics for Generative Models
ooooo

Generative Texture Models
oooooooooooo

Large-Scale GAN Training
oooooooooooooooooooooooooooo

Large-Scale Generative Modeling

Arthur Leclaire



MVA Generative Modeling
January, 23rd, 2024

Last Week

- We introduced GAN and WGAN training
- We studied a low-dimensional example explaining some instabilities...
- ... and gave some solutions to avoid them (Lipschitz penalties)
- We made some connections with optimal transport, in particular with semi-discrete WGAN

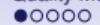
Further topics on Learning Stability

- Non-convergence of GAN/WGAN training dynamics (in simple cases) [Mescheder et al., 2018]
The authors also propose two regularizations to get convergence (continuous dynamics)
- Statistical consistency when working on a sampled version of ν
[Biau et al., 2018], [Biau et al., 2020]
- Learning multimodal distributions require generators with very large Lipschitz constants!
[Salmona et al., 2022]

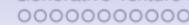
Today

- We will discuss quantitative evaluation of generative models
- We will go back to the particular case of texture generation
- We will focus on popular generative models for large-scale image generation

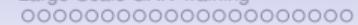
Quality Metrics for Generative Models



Generative Texture Models



Large-Scale GAN Training



Plan

Quality Metrics for Generative Models

Generative Texture Models

Large-Scale GAN Training

Quality of a Generative Model

- **Question:** How to measure that the generator covers well the training data?
 - **Main idea:** Comparing image distributions is hard...
but comparing measurements from it is easier.
 - Classification neural networks provide a set of deep non-linear features
For example, VGG19 [Simonyan and Zisserman, 2015], or Inception Networks [Szegedy et al., 2016]
 - **Measure quality of the generative model by looking at how deep statistics are preserved**
Somehow, this ensures that the database is well-covered
 - **Keep in mind that**
 - The network used to get the features must be relevant w.r.t. the generative task at play
 - Quantitative results highly depend on the network and implementation details

Inception Score \uparrow [Salimans et al., 2016]

- The inception score measures if μ generates a diverse collection of meaningful pictures
 - For an image x , Inception-v3 gives a label distribution $p(y|x)$ (discrete on $N = 1000$ labels)
 - Images containing meaningful objects have $p(y|x)$ with low entropy
 - In order to generate various images, $p(y) = \int p(y|x)\mu(dx)$ should have high entropy

The Inception Score then writes as

$$\text{IS}(\mu) = \exp\left(\int \text{KL}\left(p(y|x)|p(y)\right)\mu(dx)\right) \in [1, N]$$

It is 1 iff for a.e. x , $p(\cdot|x) = p(\cdot)$ (label distribution does not depend on x)

It is N iff for a.e. x , $p(\cdot|x)$ is concentrated on one label, and $\forall y, \int p(y|x)\mu(dx) = \frac{1}{N}$

How to compute it in practice:

- Compute an estimate $\hat{p}(y)$ of $p(y) = \int p(y|x)\mu(dx)$ by drawing samples of μ
 - Estimate $\int KL(p(y|x)|\hat{p}(y))\mu(dx)$ by drawing samples of μ

Fréchet Inception Distance (FID) ↓ [Heusel et al., 2017]

The FID measures how close are two image distributions μ, ν in terms of features distributions. It is based on the response of Inception-v3 [Szegedy et al., 2016] before last pooling layer:

$$f : \mathbb{R}^d \rightarrow \mathbb{R}^m$$

that extracts $m = 2048$ features (as a generic image summary)

NB: Images may have to be resized/normalized to fit into this network.

Algorithm to compute the FID score:

1. Draw samples (x_i) and (y_j) of $X \sim \mu$ and $Y \sim \nu$ and compute the features $(f(x_i)), (f(y_j))$
 2. Fit Gaussian distributions $\mathcal{N}(m_X, \Sigma_X)$ and $\mathcal{N}(m_Y, \Sigma_Y)$ to $(f(x_i)), (f(y_j))$ (in \mathbf{R}^{2048})
 3. Return the 2-Wasserstein distance between the Gaussian distributions,
i.e. the Fréchet distance: [Dowson and Landau, 1982]

$$W_2^2\left(\mathcal{N}(m_X, \Sigma_X), \mathcal{N}(m_Y, \Sigma_Y)\right) = \|m_X - m_Y\|_2^2 + \text{Tr}\left(\Sigma_X + \Sigma_Y - 2(\Sigma_X \Sigma_Y)^{\frac{1}{2}}\right)$$

NB: FID can be adapted to the “single-image” case: **SiFID** [Shaham et al., 2019]

SiFID compares distributions of features obtained after a convolution layer (spatially averaged)

Comments on Generative Quality

- Inception Score does not depend on the target distribution ν .
- Need to distinguish “precision/recall” for evaluating quality [Lucic et al., 2018].
“Precision” is the probability that a fake image falls within the distribution of real images.
“Recall” is the probability that a real image falls within the distribution of fake sample.
IS mainly captures precision. FID captures both precision and recall.
- The IS and FID are not enough to measure the fact that samples are photo-realistic [Barratt and Sharma, 2018],
- Other measures have been proposed better correlated with Human prediction of quality [Kolchinski et al., 2019]

Quality Metrics for Generative Models
ooooo

Generative Texture Models
●oooooooooooo

Large-Scale GAN Training
ooooooooooooooooooooooo

Plan

Quality Metrics for Generative Models

Generative Texture Models

Large-Scale GAN Training

Exemplar-based Texture Synthesis

- Exemplar texture :

$$u_0 : \Omega \rightarrow \mathbb{R}^d$$

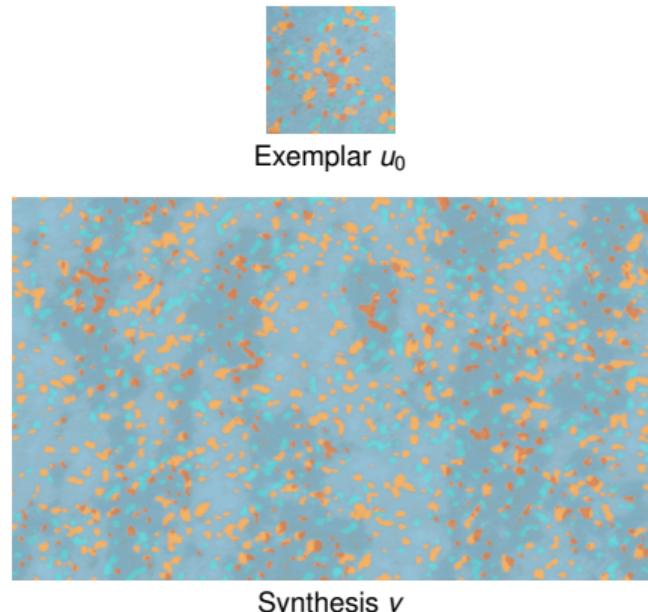
defined on a discrete rectangle $\Omega \subset \mathbb{Z}^2$.

- Texture model: stationary random field

$$V: \mathbb{Z}^2 \rightarrow \mathbb{R}^d$$

The problem can be split into

- Estimate a model V
 - Draw one (or several) samples of V



What do we want to preserve ?

- Covariance, Fourier spectrum
[Lewis, 1984], [Van Wijk, 1991], [Galerne et al., 2011], [Gilet et al., 2014]
 - Wavelet statistics
[Heeger & Bergen, 1995], [Zhu et al., 1998], [Portilla & Simoncelli, 2000],
[Tartavel et al., 2014], [Zhang & Mallat, 2017], [Bruna & Mallat, 2019]
 - Local Aspect, Patch statistics
[Efros & Leung, 1999], [Kwatra et al., 2005], [Lefebvre & Hoppe, 2005]
 - Neural Statistics
[Gatys et al., 2015], [Lu et al., 2015], [Ulyanov et al., 2016]

What do we want to preserve ?

- Covariance, Fourier spectrum
[Lewis, 1984], [Van Wijk, 1991], [Galerne et al., 2011], [Gilet et al., 2014]
 - Wavelet statistics
[Heeger & Bergen, 1995], [Zhu et al., 1998], [Portilla & Simoncelli, 2000],
[Tartavel et al., 2014], [Zhang & Mallat, 2017], [Bruna & Mallat, 2019]
 - **Local Aspect, Patch statistics**
[Efros & Leung, 1999], [Kwatra et al., 2005], [Lefebvre & Hoppe, 2005]
 - **Neural Statistics**
[Gatys et al., 2015], [Lu et al., 2015], [Ulyanov et al., 2016]

Texture Synthesis with Patch Optimal Transport

[Galerne et al., 2018]

QUESTION : How to prescribe the patch distribution at several resolutions ?

PRINCIPLE OF THE “TEXTO” MODEL:

- Initialize with a Gaussian field at coarse resolution
- At each resolution, apply a patch transport map to reimpose the exemplar patch distribution ν_s
- Upsample cleverly to go from one scale to the next

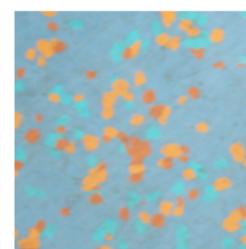


Image u^0
Patch distrib ν^0

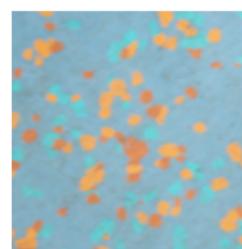


Image u^1
Patch distrib ν^1

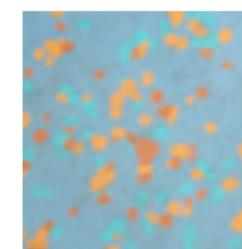


Image u^2
Patch distrib ν^2

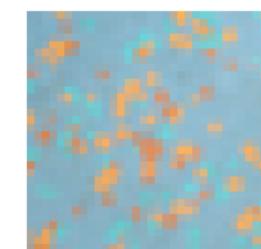


Image u^3
Patch distrib ν^3

The Texto Model

Compute exemplar $u^s : \Omega \cap (2^s \mathbb{Z}^2) \rightarrow \mathbf{R}^d$ at different scales $s = 0, \dots, S - 1$
and corresponding patch distributions ν^s

Initialize synthesis with Gaussian field U_{S-1} at the coarse scale

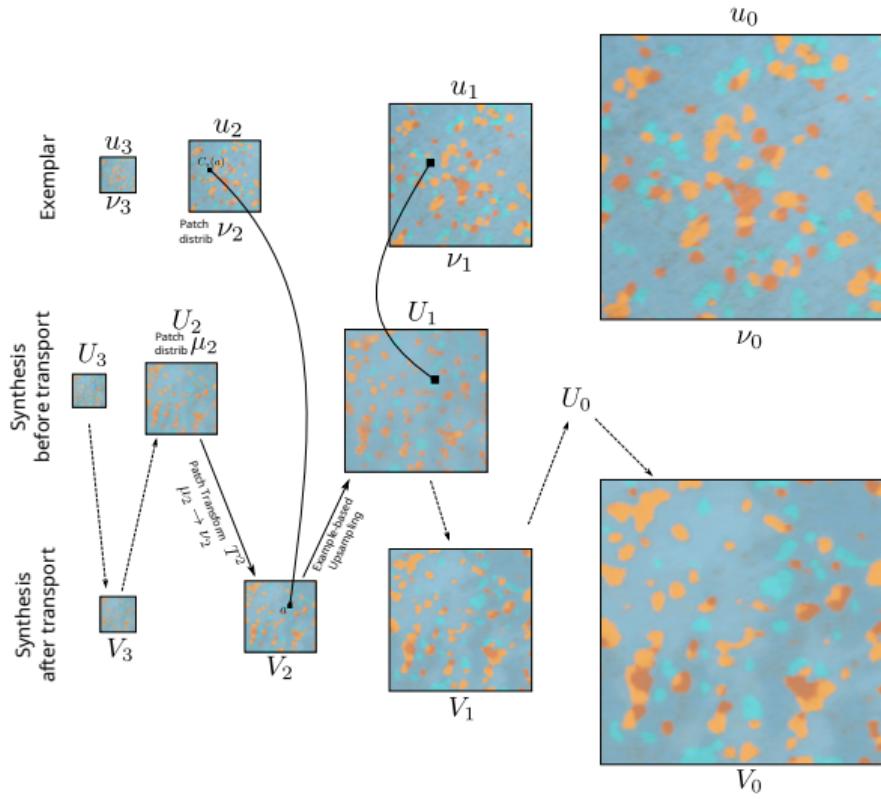
For $s = S - 1, \dots, 0$,

- Estimate the patch distribution μ_s of U_s
- Learn a patch semi-discrete OT map T_s such that $T_s \sharp \mu_s \approx \nu_s$
(Recall that T_s is a biased nearest neighbor assignment!)
- Apply T_s to all patches of U_s and recompose by averaging to an image V_s
- If $s > 0$, upsample V_s to initialize the next scale U_{s-1}
(For that, use patches at the same positions, but twice larger.)

Output: synthesis at fine scale V_0

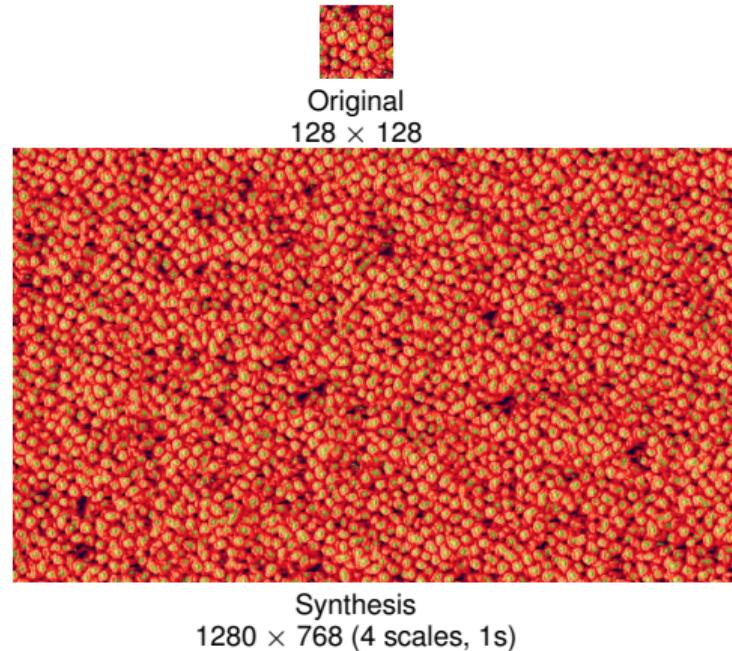
Remark: Once the model learnt, one can discard the learning steps ■ to do synthesis on-the-fly

Texto in one diagram



Texto Results

- Long-range independence property
- Patches are transformed independently
→ allows for parallel computations
- Patch OT maps can be computed offline.
→ allows for very fast synthesis
- Synthesis slightly blurry
due to patch averaging

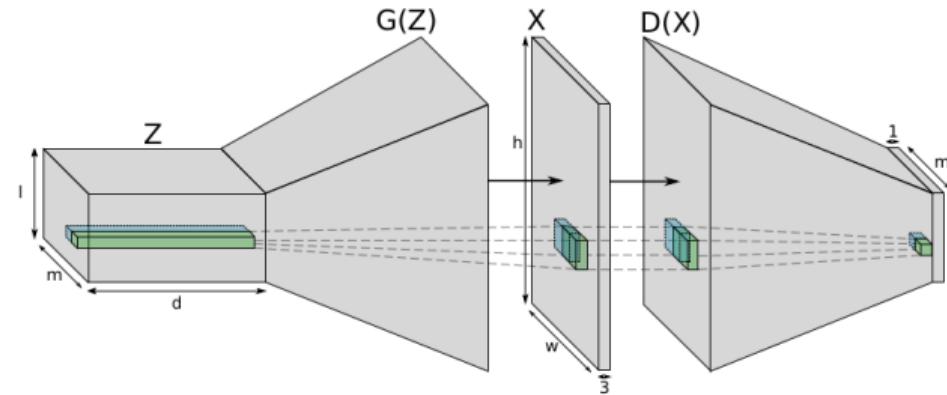


Spatial GANs [Jetchev et al., 2016], [Bergmann et al., 2017]

- Symmetric Convolutional Networks for G and D (as DCGAN, see later)
- From a $l \times m$ noise Z , $g_\theta(Z)$ generates a $h \times w$ image (in practice $l = m = 4$ and $h = w = 640$)
- Standard GAN loss (binary cross-entropy) but averaged over spatial positions (λ, μ):

$$\sum_{\lambda, \mu} \mathbb{E}[\log(1 - D_{\lambda, \mu}(g_\theta(Z)))] + \mathbb{E}[\log D_{\lambda, \mu}(Y')] \quad \text{where } Y' \text{ is a patch from } u_0$$

- PSGAN works on an augmented noise input Z , with local, global and periodic parts

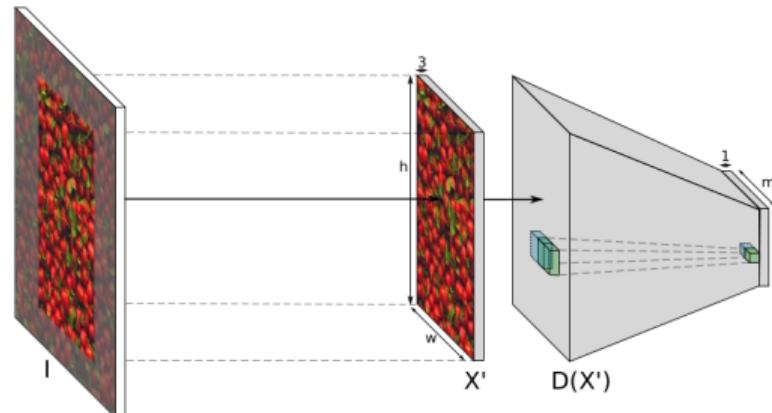


Spatial GANs [Jetchev et al., 2016], [Bergmann et al., 2017]

- Symmetric Convolutional Networks for G and D (as DCGAN, see later)
- From a $l \times m$ noise Z , $g_\theta(Z)$ generates a $h \times w$ image (in practice $l = m = 4$ and $h = w = 640$)
- Standard GAN loss (binary cross-entropy) but averaged over spatial positions (λ, μ):

$$\sum_{\lambda, \mu} \mathbb{E}[\log(1 - D_{\lambda, \mu}(g_\theta(Z)))] + \mathbb{E}[\log D_{\lambda, \mu}(Y')] \quad \text{where } Y' \text{ is a patch from } u_0$$

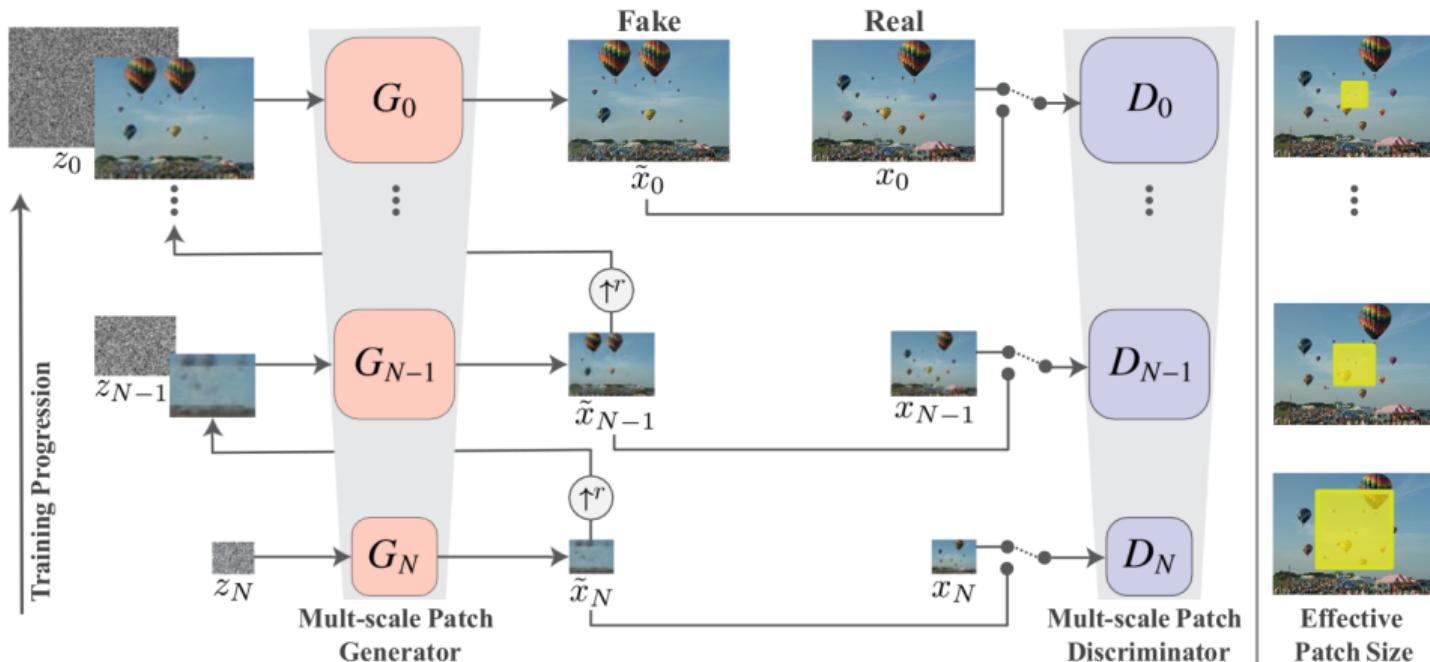
- PSGAN works on an augmented noise input Z , with local, global and periodic parts



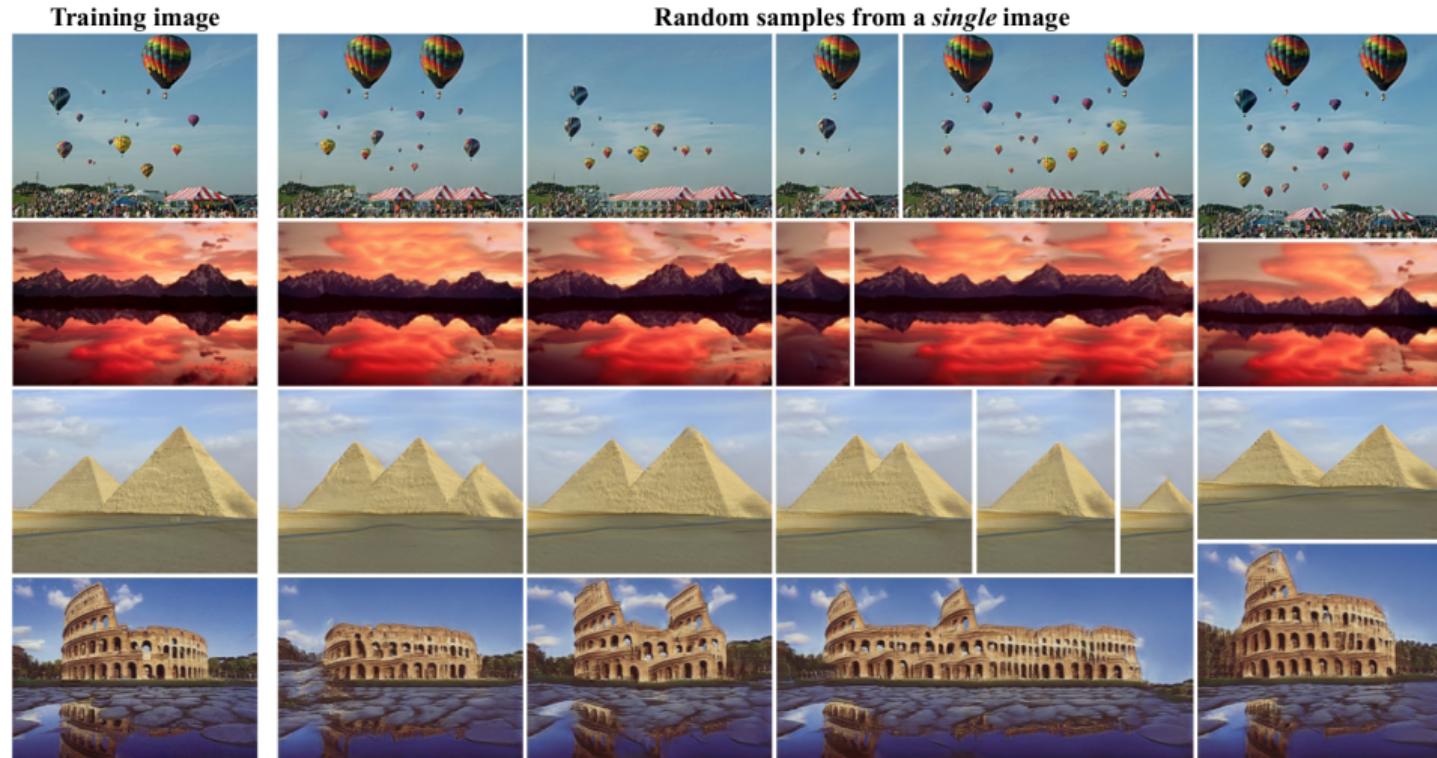
SinGAN: Learning from a Single Image [Shaham et al., 2019]

- Capture the multi-scale patch distributions of an image (possibly non-texture)
- Coarse-to-fine generator
- Patch-based dicriminator learned with WGAN-GP loss, at each scale
- Loss defined over all patches of the image, and not randomly selected patches
→ allows the network to learn boundary conditions

SinGAN: Learning from a Single Image [Shaham et al., 2019]



SinGAN: Learning from a Single Image [Shaham et al., 2019]



Generative Networks for Texture Synthesis [Houdard et al., 2023]

IDEA : Build a generative network g_θ that directly constrains features distributions where

$$\mathcal{F}_p(u) : \Omega \rightarrow \mathbf{R}^{d_p} \text{ extracts features of type } p.$$

For each feature type p , let

- $\mu_{\theta p}$: distribution of features $\mathcal{F}_p(g_\theta(Z))$
- ν_p : empirical distribution of features $\mathcal{F}_p(u_0)$

Examples:

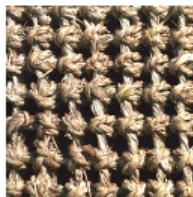
- $\mathcal{F}_p(u) : \Omega \rightarrow \mathbf{R}^{s_p \times s_p}$ extracts the $s_p \times s_p$ patches of u
- $\mathcal{F}_p(u) : \Omega_p \rightarrow \mathbf{R}^{d_p}$ extracts the response to layer p of a neural network (e.g. VGG)

Learning of GOTEX model

$$\inf_{\theta} \sum_p W(\mu_{\theta p}, \nu_p)$$

→ Alternate optimization with one dual variable ψ_p for each p

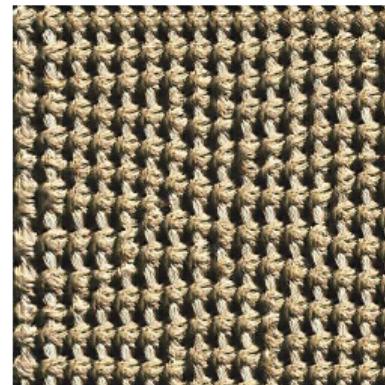
Samples of Texture Networks



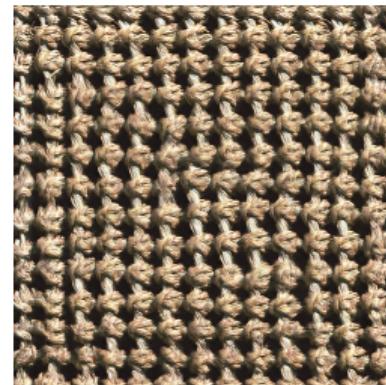
Original



GOTEX
[Houdard et al., 2023]

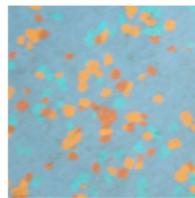


PSGAN
[Bergmann et al., 2017]

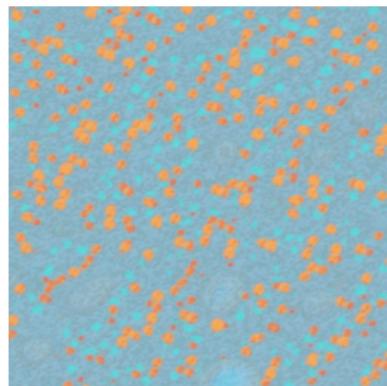


SinGAN
[Shaham et al., 2019]

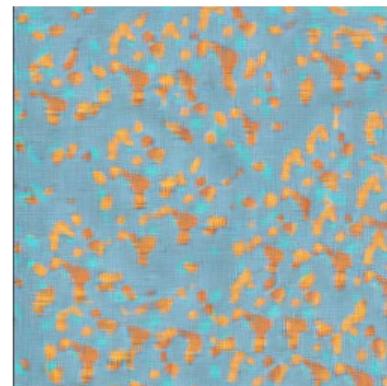
Samples of Texture Networks



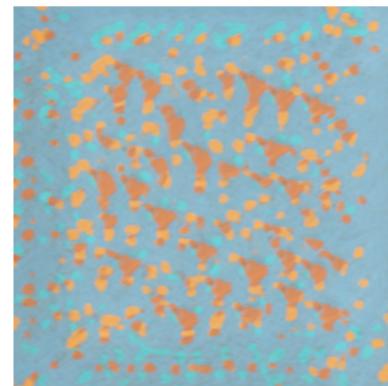
Original



GOTEX
[Houdard et al., 2023]



PSGAN
[Bergmann et al., 2017]

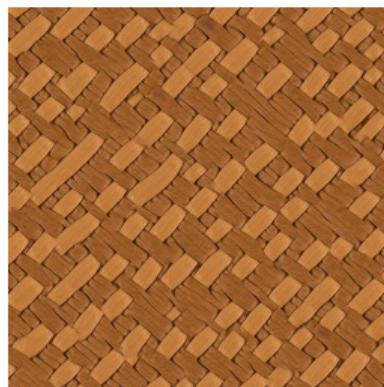


SinGAN
[Shaham et al., 2019]

Samples of Texture Networks



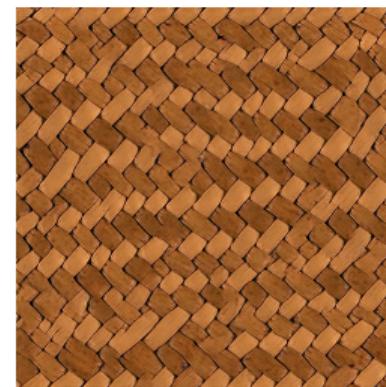
Original



GOTEX
[Houdard et al., 2023]



PSGAN
[Bergmann et al., 2017]



SinGAN
[Shaham et al., 2019]

Quality Metrics for Generative Models
ooooo

Generative Texture Models
oooooooooooo

Large-Scale GAN Training
●ooooooooooooooooooooooo

Plan

Quality Metrics for Generative Models

Generative Texture Models

Large-Scale GAN Training

Popular Image Databases

- MNIST (digits): 60000 images with 28^2 px
- Fashion-MNIST (clothes): 70000 images with 28^2 px
- CIFAR-10: 60000 images with 32^2 px
- CelebA: ≈ 200000 images with 178×278 px
- CelebA-HQ: ≈ 30000 images with 1024^2 px
- LSUN (Bedroom/Cat/Churches/...): $\approx 10^5, 10^6$ images with 256^2 px
- FFHQ (or FFHQ-U): 70000 images with 1024^2 px

Neural Network architecture

Generator and discriminator networks can have various layers:

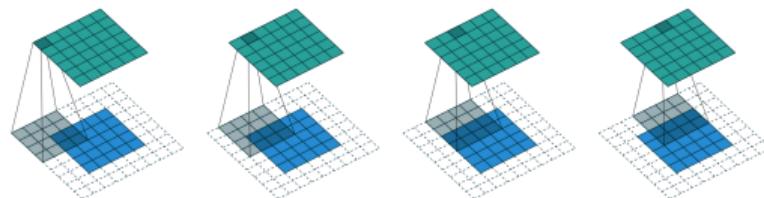
- Fully connected (FC) layers
- Upsampling (interpolation) or Subsampling (max/average pooling) layers
- Convolution/Transposed convolution (with stride), see next slide
- Activation functions: RELU, leakyRELU, sigmoid, tanh, etc
- BatchNorm
- ...

Convolution and Transposed convolution

The convolution of u, v is defined by

$$w * u(i) = \sum_j w(j)u(i-j),$$

where $u(j) \in \mathbf{R}^d$, $w(j) \in \mathbf{R}^{d' \times d}$.



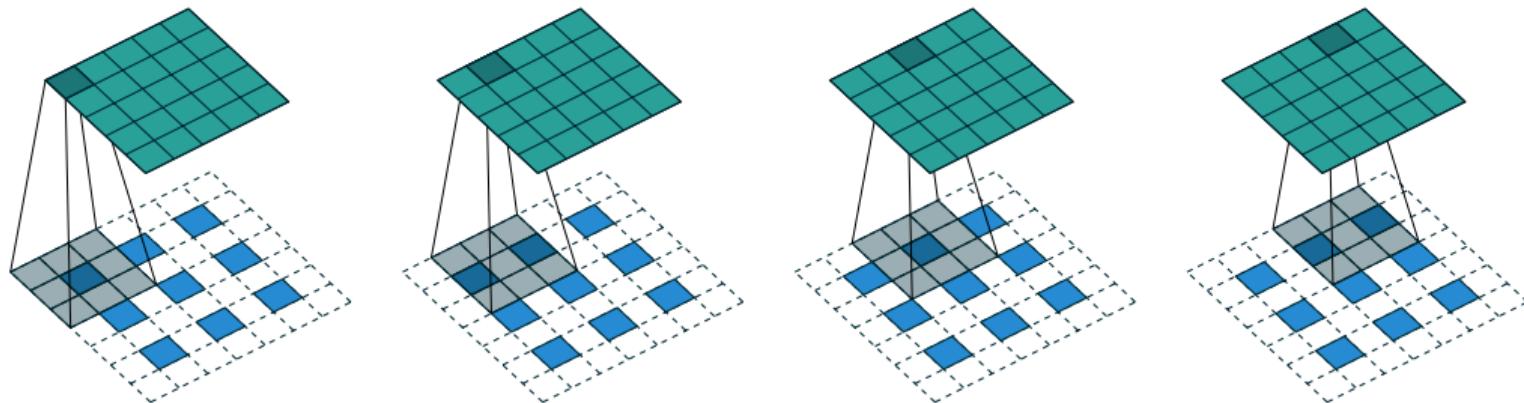
Notice that

- The transpose of a convolution with a $k \times k$ kernel is a convolution with a $k \times k$ kernel
- The transpose of a border crop is zero-padding the borders.
- The transpose of a crude subsampling is zero-inserting.

Fractionally strided convolutions:

- This is the transpose operator of convolution+subsampling (convolution with stride).
- Called `ConvTranspose2d` in PyTorch

One Example from [Dumoulin and Visin, 2016]



"The transpose of convolving a 3×3 kernel over a 5×5 input padded with a 1×1 border of zeros using 2×2 strides (i.e., $i = 5$, $k = 3$, $s = 2$ and $p = 1$). It is equivalent to convolving a 3×3 kernel over a 3×3 input (with 1 zero inserted between inputs) padded with a 1×1 border of zeros using unit strides (i.e., $i' = 3$, $\tilde{i}' = 5$, $k' = k$, $s' = 1$ and $p' = 1$)."

BatchNorm layer

Principle of BatchNormalization: for any batch $(x_i)_{i \in B}$ of a K -dimensional feature, transform

$$\forall k = 1, \dots, K, \quad y_i^{(k)} = \gamma^{(k)} \hat{x}_i^{(k)} + \beta^{(k)} \quad \text{with} \quad \hat{x}_i^{(k)} = \frac{x_i^{(k)} - m_i^{(k)}}{\sqrt{(\sigma_i^{(k)})^2 + \varepsilon}}$$

where $m_i^{(k)}, \sigma_i^{(k)}$ are mean and std of the k -feature over this batch,
and where $\gamma^{(k)}, \beta^{(k)}$ are trainable parameters.

At inference: normalization is done with $m^{(k)}, \sigma^{(k)}, \gamma^{(k)}, \beta^{(k)}$ learned during training.
Switch to inference mode with `model.eval()`.

Convolutional GAN

[Radford et al., 2016]

Important principles of the construction:

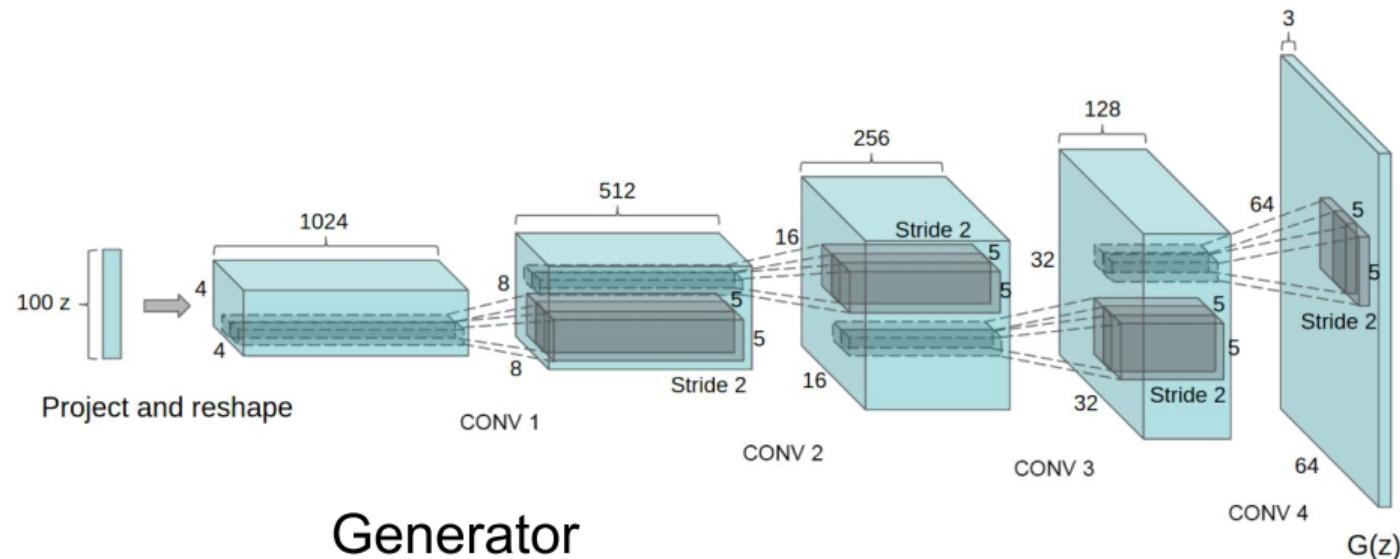
- “All convolutional”: remove max pooling layers, and learn downsampling instead
- Eliminate Fully-Connected Layers
- Batch Normalization to stabilize learning (except on generator output, and discriminator input)
- ReLU activations for the generator
- LeakyReLU activations for the discriminator

Generator: upsampling network with **fractionally strided convolutions** (i.e. the transpose operator of convolution+subsampling , called `ConvTranspose2d` in PyTorch)

Discriminator: convolutional network with strided convolutions

DCGAN Architecture

[Radford et al., 2016]



Generator

Image Generation with DCGAN [Radford et al., 2016]



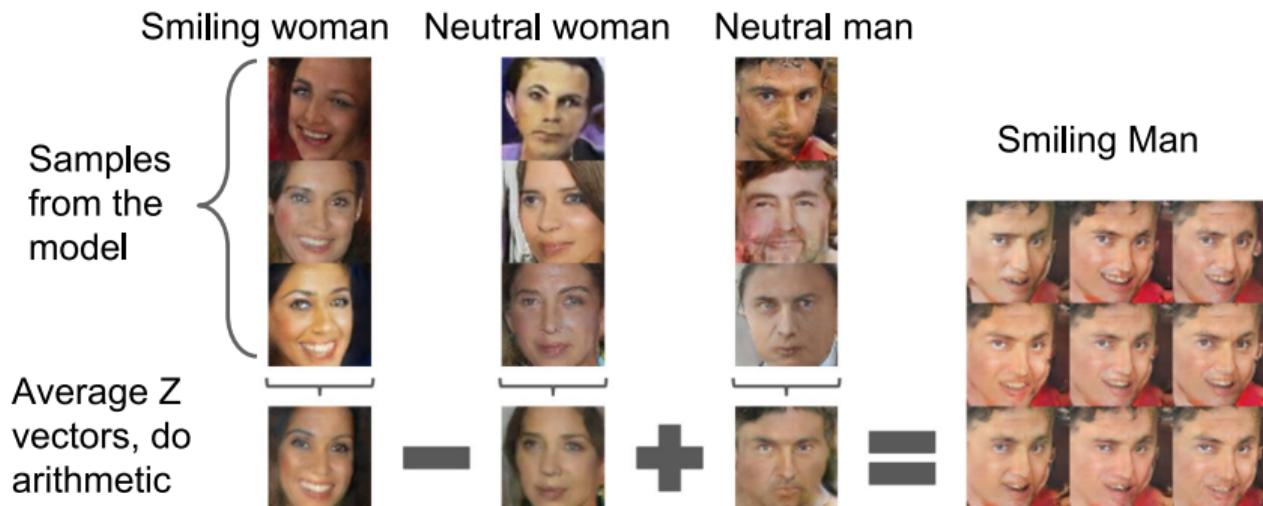
Generations of realistic bedrooms pictures, from randomly generated latent variables.

Image Interpolation with DCGAN [Radford et al., 2016]



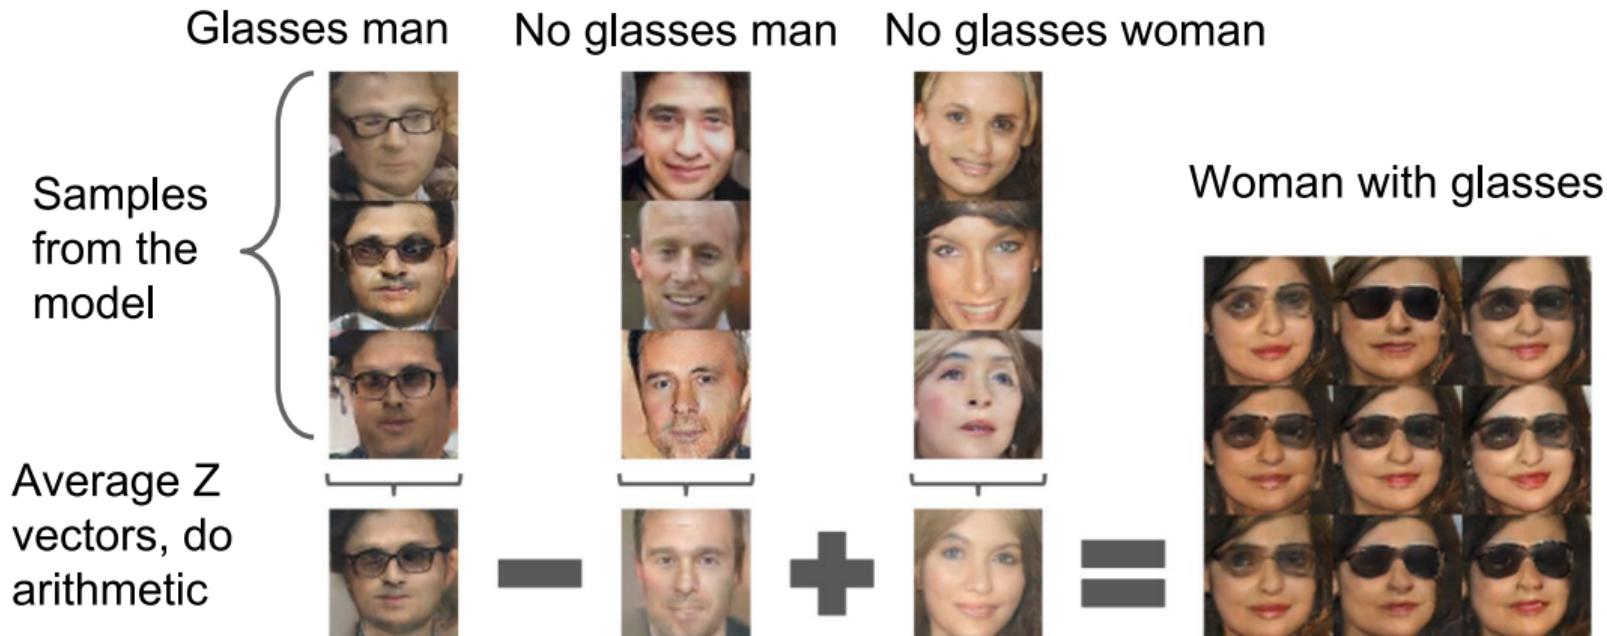
Interpolation in between points in latent space.

Arithmetic with DCGAN [Radford et al., 2016]



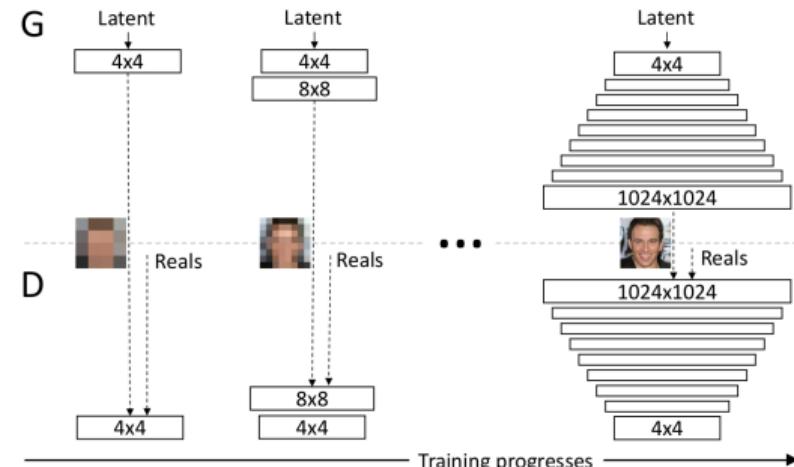
- Average latent vector of several samples
- After arithmetic, add a small random perturbation to get similar samples

Arithmetic with DCGAN [Radford et al., 2016]



Progressive Growing of GANs [Karras et al., 2018]

- Progressive Multiresolution Training
- Mirror architectures for G and D
- Simple upsampling/downsampling
nearest neighbor upsampling;
average pooling downsampling
- Minibatch statistics layer at the end of D
- Pixelwise feature normalization
- Training with WGAN-GP



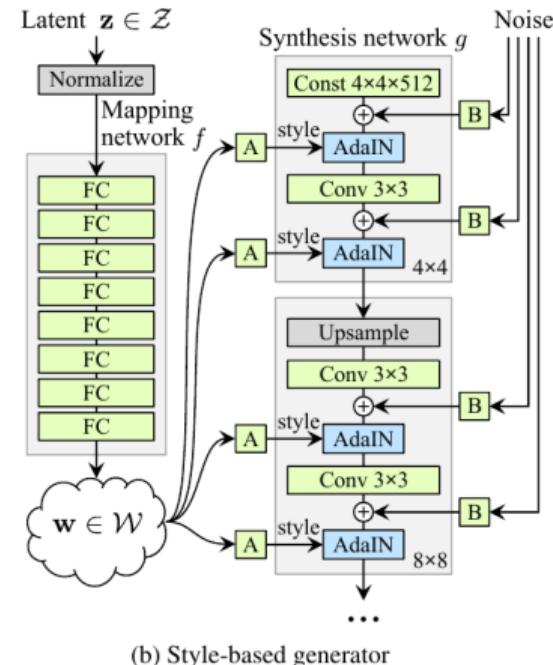
StyleGAN [Karras et al., 2019]

- “separation of high-level features (pose, identity) from stochastic variation (freckles, hair)”
- Embed input latent code z into an intermediate latent space w with a multilayer perceptron (8 FC layers)
- Spatially invariant style vector $y = (y_s, y_b)$ for each feature map, obtained from w
- AdaIN: Adaptive Instance Normalization

$$\text{AdaIN}(x_i, y) = y_{s,i} \frac{x_i - \mu(x_i)}{\sigma(x_i)} + y_{b,i}$$

where the feature map x_i is normalized separately

- Style mixing (playing with two latent codes w_1, w_2)



StyleGAN [Karras et al., 2019]

- “separation of high-level features (pose, identity) from stochastic variation (freckles, hair)”
- Embed input latent code z into an intermediate latent space w with a multilayer perceptron (8 FC layers)
- Spatially invariant style vector $y = (y_s, y_b)$ for each feature map, obtained from w
- AdaIN: Adaptive Instance Normalization

$$\text{AdaIN}(x_i, y) = y_{s,i} \frac{x_i - \mu(x_i)}{\sigma(x_i)} + y_{b,i}$$

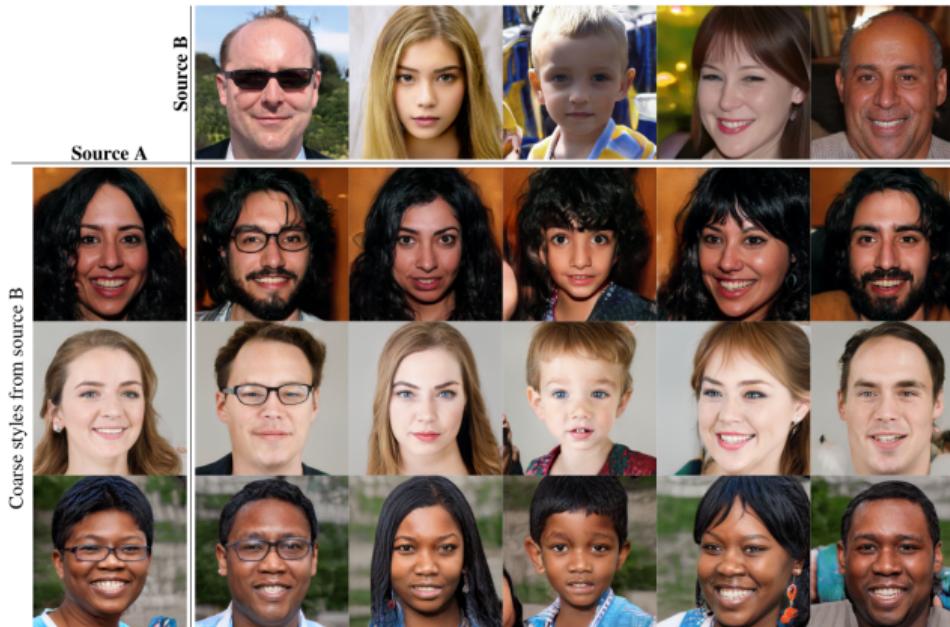
where the feature map x_i is normalized separately

- Style mixing (playing with two latent codes w_1, w_2)



StyleGAN [Karras et al., 2019]

StyleGAN allows for style mixing at different scales (by using the corresponding subpart of w).



StyleGAN2 [Karras et al., 2020]

- AdaIN causes droplet artifacts in StyleGAN
→ Weight modulation/demodulation instead of AdaIN
- Path length regularization:
fixed step-size in w results in fixed magnitude change in imag
- Residual connections with downsampling in D
- Skip connections in G
- No progressive growing (which leads to *phase artifacts*)



StyleGAN vs StyleGAN2



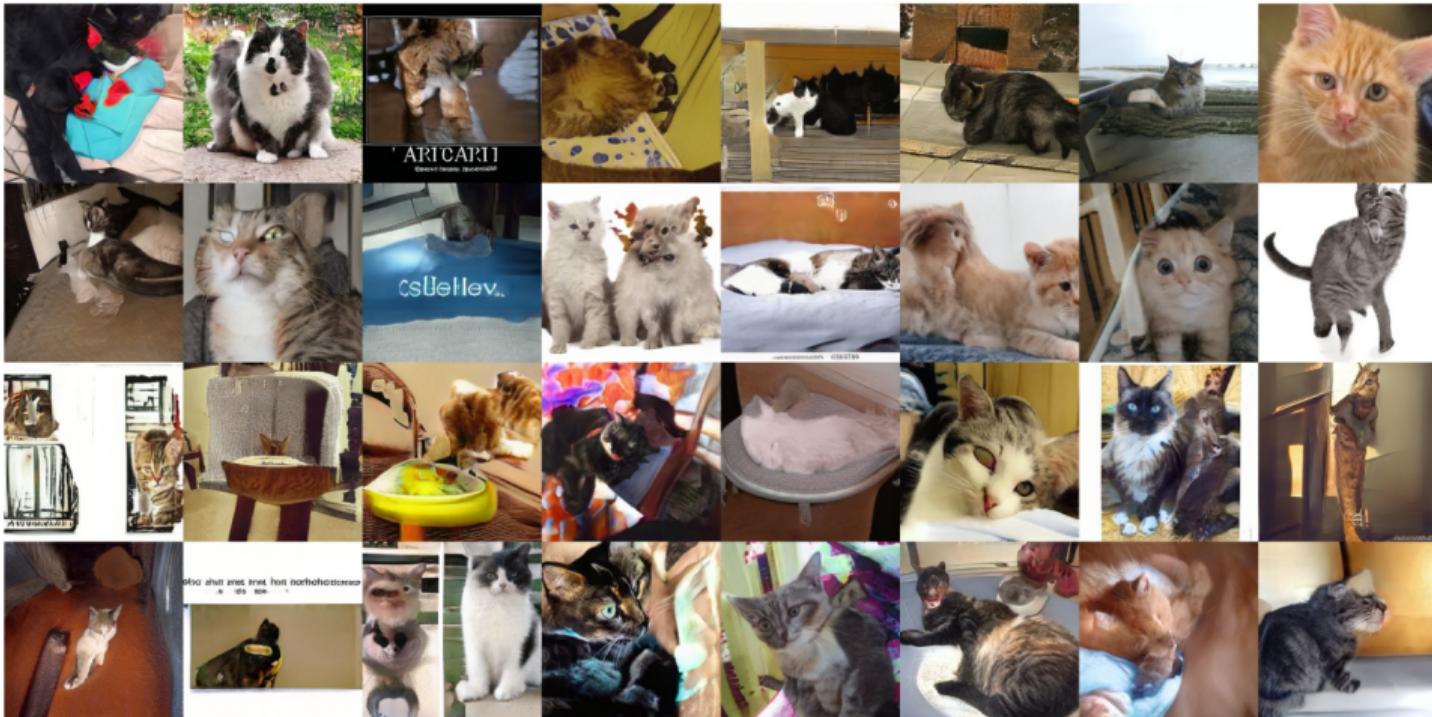
First row: real images
Second row: samples of StyleGAN after projection on the latent code

StyleGAN vs StyleGAN2



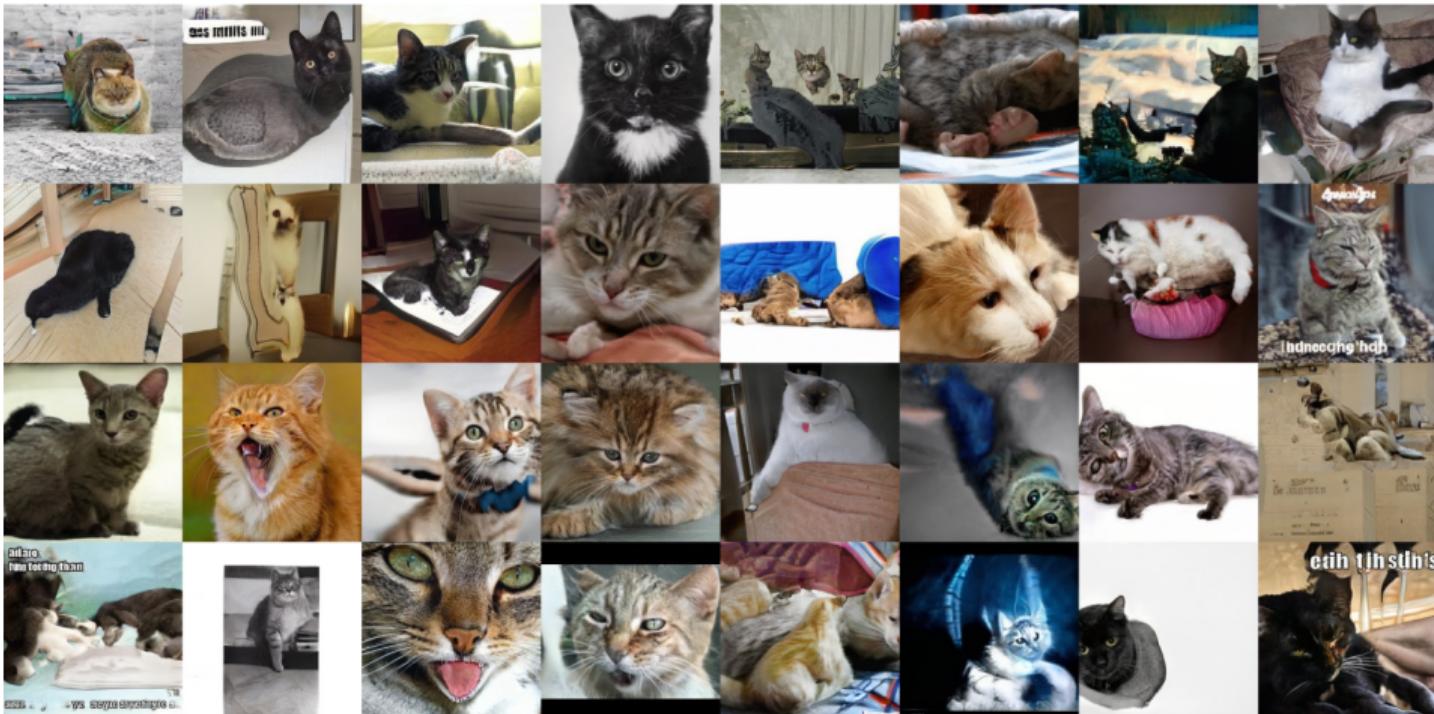
First row: real images
Second row: samples of StyleGAN2 after projection on the latent code

The Cat Challenge...



Samples of StyleGAN2-Model1 trained on LSUN Cat

The Cat Challenge...



Samples of StyleGAN2-Model2 trained on LSUN Cat

StyleGAN3 aka Alias-free GANs

- Aliasing artifacts present in some GANs results due to:
 - non-ideal upsampling
 - pointwise activations
- Enforce continuous equivariance to sub-pixel translation (Shannon is back...)
- Also, ensure that no aliasing appears through the network:
 - use band-limited filters
 - use low-pass filters when needed

StyleGAN3 aka Alias-free GANs

Are GANs created equal?

[Lucic et al., 2018]

Many variants of GAN training exist, with various architectures and more or less stable training.

- Regarding quality of generated images, may GAN variants perform similarly.
Lucic et al. proposed a large comparison framework, with a budget for hyperparameter tuning, and by averaging over several random seeds.
- “WGANS work because they fail” [Stanczuk et al., 2021], [Mallasto et al., 2019]
The dual training in WGAN-GP does not approximate the Wasserstein distance correctly.
But estimating it more precisely (e.g. semi-discrete WGAN) often leads to blurrier samples.
→ The quality of a generative network relies on good features learned by the discriminator.

Take-home Messages

- FID score gives a reasonable/simple way to measure the quality of a generative model... but it does not suffice to judge photo-realism of the samples
- We discussed several architectures for texture/image generation
- Large-scale synthesis benefits from architectures adapted for multi-resolution synthesis.
- Recent generative models crucially rely on
 - several tricks for training or designing the architecture
 - very long training of models with an extremely large number of parameters

THANK YOU FOR YOUR ATTENTION!

References I

-  Barratt, S. and Sharma, R. (2018).
A note on the inception score.
arXiv preprint arXiv:1801.01973.
-  Bergmann, U., Jetchev, N., and Vollgraf, R. (2017).
Learning texture manifolds with the periodic spatial gan.
In *Proceedings of the 34th International Conference on Machine Learning-Volume 70*, pages 469–477. JMLR. org.
-  Biau, G., Cadre, B., Sangnier, M., and Tanielian, U. (2018).
Some theoretical properties of gans.
arXiv preprint arXiv:1803.07819.
-  Biau, G., Sangnier, M., and Tanielian, U. (2020).
Some theoretical insights into wasserstein gans.
-  Dowson, D. C. and Landau, B. V. (1982).
The fréchet distance between multivariate normal distributions.
Journal of Multivariate Analysis, 12(3):450–455.



References II

-  [Dumoulin, V. and Visin, F. \(2016\).](#)
A guide to convolution arithmetic for deep learning.
ArXiv e-prints.
-  [Galerne, B., Leclaire, A., and Rabin, J. \(2018\).](#)
A texture synthesis model based on semi-discrete optimal transport in patch space.
SIAM Journal on Imaging Sciences, 11(4):2456–2493.
-  [Heusel, M., Ramsauer, H., Unterthiner, T., Nessler, B., and Hochreiter, S. \(2017\).](#)
Gans trained by a two time-scale update rule converge to a local Nash equilibrium.
In Guyon, I., Luxburg, U. V., Bengio, S., Wallach, H., Fergus, R., Vishwanathan, S., and Garnett, R., editors, *Advances in Neural Information Processing Systems*, volume 30. Curran Associates, Inc.
-  [Houdard, A., Leclaire, A., Papadakis, N., and Rabin, J. \(2023\).](#)
A generative model for texture synthesis based on optimal transport between feature distributions.
Journal of Mathematical Imaging and Vision, 65(1):4–28.



References III

-  Jetchev, N., Bergmann, U., and Vollgraf, R. (2016).
Texture synthesis with spatial generative adversarial networks.
-  Karras, T., Aila, T., Laine, S., and Lehtinen, J. (2018).
Progressive Growing of GANs for Improved Quality, Stability, and Variation.
In *Proceedings of International Conference on Learning Representations*.
-  Karras, T., Laine, S., and Aila, T. (2019).
A style-based generator architecture for generative adversarial networks.
In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 4401–4410.
-  Karras, T., Laine, S., Aittala, M., Hellsten, J., Lehtinen, J., and Aila, T. (2020).
Analyzing and improving the image quality of StyleGAN.
In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pages 8110–8119.
-  Kolchinski, Y. A., Zhou, S., Zhao, S., Gordon, M. L., and Ermon, S. (2019).
Approximating human judgment of generated image quality.
CoRR, abs/1912.12121.



References IV

-  **Lucic, M., Kurach, K., Michalski, M., Gelly, S., and Bousquet, O. (2018).**
Are GANs created equal? A large-scale study.
In Bengio, S., Wallach, H., Larochelle, H., Grauman, K., Cesa-Bianchi, N., and Garnett, R., editors, *Advances in Neural Information Processing Systems*, volume 31. Curran Associates, Inc.
-  **Mallasto, A., Montúfar, G., and Gerolin, A. (2019).**
How well do WGANs estimate the Wasserstein metric?
arXiv preprint arXiv:1910.03875.
-  **Mescheder, L., Geiger, A., and Nowozin, S. (2018).**
Which training methods for gans do actually converge?
arXiv preprint arXiv:1801.04406.
-  **Radford, A., Metz, L., and Chintala, S. (2016).**
Unsupervised representation learning with deep convolutional generative adversarial networks.
In Bengio, Y. and LeCun, Y., editors, *Proceedings of ICLR*.



References V

-  Salimans, T., Goodfellow, I., Zaremba, W., Cheung, V., Radford, A., and Chen, X. (2016). Improved techniques for training gans.
Advances in neural information processing systems, 29.
-  Salmona, A., De Bortoli, V., Delon, J., and Desolneux, A. (2022). Can push-forward generative models fit multimodal distributions?
Advances in Neural Information Processing Systems, 35:10766–10779.
-  Shaham, T. R., Dekel, T., and Michaeli, T. (2019). SinGAN: Learning a Generative Model from a Single Natural Image.
In *Proceedings of the IEEE International Conference on Computer Vision*, pages 4570–4580.
-  Simonyan, K. and Zisserman, A. (2015). Very deep convolutional networks for large-scale image recognition.
In Bengio, Y. and LeCun, Y., editors, *Proceedings of the International Conference on Learning Representations*.
-  Stanczuk, J., Etmann, C., Kreusser, L. M., and Schönlieb, C.-B. (2021). Wasserstein GANs work because they fail (to approximate the Wasserstein distance).
arXiv preprint arXiv:2103.01678.

References VI

- 
- Szegedy, C., Vanhoucke, V., Ioffe, S., Shlens, J., and Wojna, Z. (2016). Rethinking the inception architecture for computer vision. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*.